

# STATS 230: Computational Statistics

## Numerical linear algebra

Babak Shahbaba

Department of Statistics, UCI

# Overview

- We are interested in solving equations  $Ax = b$
- First, I will go over some basic concepts in linear algebra
- Next, I will talk about numerical linear algebra for developing fast computational methods
- Finally, I will discuss these methods in the context of linear regression models
- The review of linear algebra, algorithms, and most of examples presented here are mainly based on Strang (2012)
- For more details, refer to Strang (2012), Boyd and Vandenberghe (2004), and Thisted (1988)

# Some important concepts in linear algebra

# Four fundamental spaces

- Conceptually,  $Ax = b$  should be interpreted as “ $A$  acts on  $x$  to produce  $b$ ”.
- Further, we can think of  $Ax$  as a linear combinations of the columns of  $A$ :  $x_1a_1 + x_2a_2, \dots, x_na_n$ , where  $a_1, a_2, \dots, a_n$  are the columns of  $A$ .
- All possible combinations of the columns form the columns space  $C(A)$
- $Ax = b$  is solvable if  $b \in C(A)$
- The null space  $N(A)$  on the other hand includes all vectors  $x$  such that  $Ax = 0$
- For full column rank matrices,  $N(A)$  contains only zero:  $x = (0, \dots, 0)$ .
- When  $A$  is  $m$  by  $n$ , then  $C(A)$  is a subspace of  $R^m$  and  $N(A)$  is a subspace of  $R^n$ .

# Four fundamental spaces

- We can also talk about two other subspaces:  $C(A^\top)$ , which is also called the row space, and  $N(A^\top)$ .
- Together,  $C(A)$ ,  $N(A)$ ,  $C(A^\top)$ , and  $N(A^\top)$  create four fundamental subspaces, which are very important in linear algebra.
- Recall that the dimension,  $r$ , of a space is the number of independent vectors.
- We can show that
  - ▶ The column space  $C(A)$  in  $R^m$  and the row space  $C(A^\top)$  in  $R^n$  have the same dimension  $r$ .
  - ▶ The null spaces  $N(A)$  and  $N(A^\top)$  have dimensions  $n - r$  and  $m - r$  respectively.

# Four fundamental spaces

- Recall that if  $V$  is a subspace of  $R^n$ , its orthogonal complement is

$$V^\perp = \{x | z^\top x = 0, \forall z \in V\}$$

then, we can write each vector in  $R^n$  as a sum of two vectors from  $V$  and  $V^\perp$

$$R^n = V \overset{\perp}{\oplus} V^\perp$$

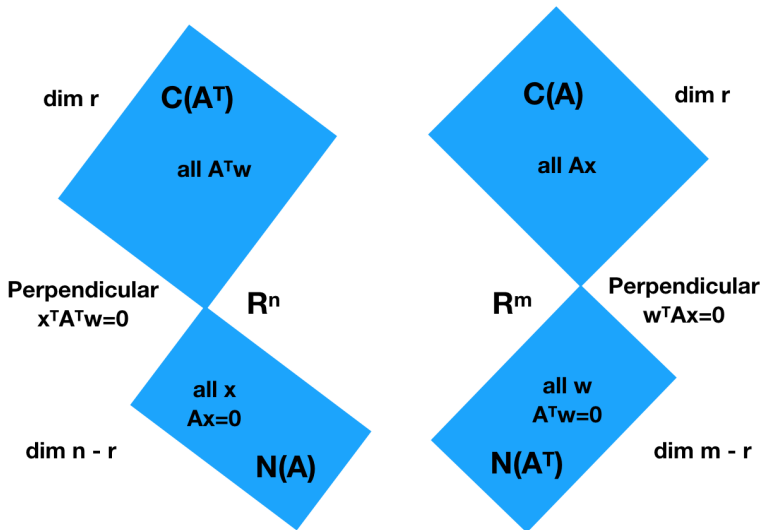
- Given  $A_{m \times n}$ , for the four fundamental subspaces mentioned above we have

$$N(A^\top) \overset{\perp}{\oplus} C(A) = R^m$$

$$N(A) \overset{\perp}{\oplus} C(A^\top) = R^n$$

# Four fundamental spaces

- Schematically, the four subspaces can be presented as follows (Strang, 2012).



# Basis

- A full set of independent vectors form a basis for a space
- Each vector in the space can be presented as a unique combination of these basis vectors
- Possible choices:
  - ▶ Standard basis: columns of the identity matrix
  - ▶ General basis: columns of any invertible matrix
  - ▶ Orthogonal basis: columns of any orthogonal matrix



# Orthogonal matrices

- Note that for a matrix  $Q$  with orthonormal columns,  $q_1, \dots, q_n$ , we have

$$q_i^\top q_j = \begin{cases} 0 & i \neq j \\ 1 & i = j \end{cases} \quad Q^\top Q = I$$

- If  $Q$  is a square matrix, it is called an *orthogonal matrix* and  $Q^\top = Q^{-1}$
- Multiplying a vector by  $Q$  doesn't change its length:

$$\|Qx\|^2 = x^\top Q^\top Qx = x^\top x = \|x\|^2$$

# Eigenvalues and eigenvectors

- When  $A$  acts on  $x$  (i.e.,  $Ax$ ), it almost always changes the direction of  $x$
- For some special vectors,  $Ax = \lambda x$  so  $x$  either stretches, shrinks, reverses directions, or stays unchanged
- Then, we say  $x$  is an *eigenvector* for  $A$  and  $\lambda$  is its corresponding *eigenvalue*
- One possible way (good for low-dimensional problems) to find  $x$  and  $\lambda$  is through solving

$$(A - \lambda I)x = 0$$

since  $(A - \lambda I)$  needs to be singular so its null space includes  $x \neq 0$ , we have

$$\det(A - \lambda I) = 0$$

- This is called the characteristic equation, which involves  $\lambda$  only (not  $x$ ).

# Eigenvalues and eigenvectors

- After we find  $\lambda$ 's, we can find the corresponding eigenvectors. We can also calculate

$$\det(A) = \prod_{i=1}^n \lambda_i \quad \text{trace}(A) = \sum_{i=1}^n \lambda_i$$

- If  $A$  is triangular, then its eigenvalues are given by its diagonal elements
- For square matrices, the eigenvalues of  $A^2$  are  $\lambda_1^2, \dots, \lambda_n^2$

$$Ax = \lambda x \Rightarrow A^2x = \lambda Ax = \lambda^2 x; \quad \text{In general, } A^k x = \lambda^k x$$

and the eigenvalues of  $A^{-1}$  are  $1/\lambda_1, \dots, 1/\lambda_n$

$$Ax = \lambda x; A^{-1}Ax = \lambda A^{-1}x; A^{-1}x = \frac{1}{\lambda}x$$

- Note that the eigenvectors remain the same

# Diagonalization

- Suppose  $A_{n \times n}$  has  $n$  independent eigenvectors,  $x_1, \dots, x_n$ , which are the columns of an eigenvector matrix  $S$
- The corresponding eigenvalues form a diagonal matrix  $\Lambda = \text{diag}(\lambda_1, \dots, \lambda_n)$
- Then, we can write them in a matrix form  $AS = S\Lambda$ , from which we get

$$\begin{aligned} S^{-1}AS &= \Lambda \\ A &= S\Lambda S^{-1} \end{aligned}$$

- Symmetric matrices have real eigenvalues and orthogonal eigenvector matrix  $Q$  where,

$$q_i^\top q_j = \begin{cases} 0 & i \neq j \\ 1 & i = j \end{cases}$$

therefore,  $Q^\top Q = I$ ;  $Q^\top = Q^{-1}$ , and  $A = Q\Lambda Q^\top$

# Positive definiteness

- A square matrix,  $S$ , is positive definite if  $u^T S u > 0$ ,  $\forall u \neq 0$
- You can think of  $\frac{1}{2} u^T S u$  as the *energy* function of a system. We can find its minimum by setting the gradient (with respect to  $u$ ) to zero:  $Su = 0$ , and show that this point is in fact the minimum if the Hessian (second derivatives)  $S$  is positive definite.
- For positive definite matrices, all eigenvalues are positive
- $S = A^T A$  is symmetric and positive definite (when  $A$  has independent columns so the only solution to  $Ax = 0$  is the zero vector) or at least semidefinite (when  $Ax = 0$  has nonzero solutions)
- If  $S_1$  and  $S_2$  are positive definite, then  $S_1 + S_2$  is also positive definite

# Singular vs. nonsingular

- As mentioned above, we are primarily interested in solving equations  $Ax = b$
- If possible, we could solve the above equation as  $x = A^{-1}b$
- We can express whether this is possible or not in different ways

## Nonsingular

$A$  is invertible

$Ax = b$  has one solution:  $A^{-1}b$

$Ax = 0$  has one solution:  $x = 0$

The columns are independent

The rows are independent

The column space is  $\mathbb{R}^n$

The row space is  $\mathbb{R}^n$

$A$  has full rank

$A$  has  $n$  positive singular values

$A^T A$  is symmetric positive definite

All eigenvalues of  $A$  are nonzero

The determinant is nonzero

## Singular

$A$  is not invertible

$Ax = b$  has no solution or infinitely many solutions

$Ax = 0$  has many solutions

The columns are dependent

The rows are dependent

The column space has  $\dim r < n$

The row space has  $\dim r < n$

$A$  has  $r < n$  rank

$A$  has  $r < n$  singular values

$A^T A$  is only semidefinite

Zero is an eigenvalue of  $A$

The determinant is zero

# Numerical linear algebra

# Flops

- In general, solving  $Ax = b$  is difficult for big matrices
- The computational cost is lower when working with structured matrices: symmetric, triangular, orthonormal, sparse, diagonal
- The computational cost of algorithms in numerical linear algebra is commonly measured by the total number of floating-point operations (flops)
- A flop is one addition, subtraction, multiplication, or division of two floating-point numbers
- To evaluate the computational cost of an algorithm, we count the total number of flops as a function of the dimensions of matrices and vectors
- This is usually a polynomial function, and we typically focus on the dominant (higher order) terms by using the big-O notation:  $\mathcal{O}$ .



Operation	Cost
Inner product	$\mathcal{O}(n)$
$A_{m \times n} x$	$\mathcal{O}(mn)$
$A_{m \times n} N_{n \times p}$	$\mathcal{O}(mnp)$
Solving $Ax = b$ ; $A$ is dense	$\mathcal{O}(n^3)$
Solving $Ax = b$ ; $A$ is orthogonal	$\mathcal{O}(n^2)$
Solving $Ax = b$ ; $A$ is triangular	$\mathcal{O}(n^2)$
Solving $Ax = b$ ; $A$ is banded with bandwidth $k$	$\mathcal{O}(k^2 n)$
Solving $Ax = b$ ; $A$ is diagonal	$\mathcal{O}(n)$

# Factorization

- As mentioned earlier, it is easier and more computationally efficient to work with structured matrices.
- We now discuss three types of factorizations to generate such matrices
  - ▶  $A = LU$  = lower triangular  $\times$  upper triangular
  - ▶  $A = QR$  = Orthonormal columns  $\times$  upper triangular
  - ▶  $A = U\Sigma V^T$ , where  $U$  and  $V$  are orthonormal and  $\Sigma$  is a diagonal matrix

# LU Factorization

- To solve  $Ax = b$ , we could use a set of forward elimination operations to change  $A$  to an upper triangular matrix  $U$ ,

$$U = \begin{pmatrix} u_{11} & u_{12} & \dots & u_{1n} \\ 0 & u_{22} & \dots & \vdots \\ \vdots & \vdots & \ddots & \vdots \\ 0 & \dots & \dots & u_{nn} \end{pmatrix}$$

- Changing the problem to  $Ux = c$ , we solve the system by backward substitution starting from the last equation.
- It turns out a lower triangular matrix,  $L$ , can take  $U$  back to  $A$  such that

$$A = LU$$

- This provides a factorization for  $A$

# LU Factorization

- Sometimes, we need to apply some row permutations to  $A$  before factorizing it,

$$PA = LU$$

note that permutations cost zero flops.

- The factorization usually cost  $(2/3)n^3$  flops, but it could be much lower for sparse matrices
- We can then solve the equation  $PAx = LUx = Pb$  as follows
  - ▶ 1)  $z_1 = Pb$ ; zero flops
  - ▶ 2)  $Lz_2 = z_1$ ;  $n^2$  flops
  - ▶ 3)  $Ux = z_2$ ;  $n^2$  flops

# Cholesky Factorization

- For symmetric matrices, we obtain symmetric factorizations

$$A = LDL^{\top}$$

where  $D$  is a diagonal matrix so

$$A^{\top} = (L^{\top})^{\top}DL^{\top} = LDL^{\top}$$

- Alternatively, we can write this as

$$A = L\sqrt{D}\sqrt{D}L^{\top} = L^*L^{*\top}$$

- This factorization costs  $(1/3)n^3$  flops

# Orthogonalization: $QR$ factorization

- A common factorization is  $A = QR$ , where  $R$  is an upper triangular matrix and  $Q$  has orthonormal columns,  $Q^T Q = I$
- Then, instead of columns  $a_1, \dots, a_n$  as basis, we would use  $q_1, \dots, q_n$
- Two common methods for  $QR$  factorization are Gram-Schmidt and Householder (discussed later)
- Using this factorization, operations involving  $A^T A$  simplify to  $Q^T Q = I$
- Additionally, because of  $Q$ 's *stability*, we can avoid overflow and underflow
- Finally, since multiplying by  $Q$  does not change the size, small numerical errors  $\Delta b$  will not generate very large errors in  $x$ :

$$Q(\Delta x) = \Delta b \Rightarrow \|\Delta x\| = \|\Delta b\|$$

# Gram-Schmidt

- We start with setting  $q_1 = a_1 / \|a_1\|$  and since  $a_1 = r_{11}q_1$ , we have  $r_{11} = \|a_1\|$
- To find  $q_2$ , we subtract from  $a_2$  its component in  $q_1$  direction and then normalize:

$$w_2 = a_2 - (q_1^\top a_2)q_1; \quad q_2 = w_2 / \|w_2\|$$

- At step  $k$ , we subtract from  $a_k$  its projection on  $q_1, \dots, q_{k-1}$  and then normalize
- Finally, we will have  $A_{m \times n} = Q_{m \times n} \times R_{n \times n}$
- In contrast, the Householder algorithm (discussed later) creates  $A_{m \times n} = Q_{m \times m} \times R_{m \times n}$

# Gram-Schmidt

Initialize  $Q_{m \times n} = 0$ ,  $R_{n \times n} = 0$ , and  $v = A[:, 1]$

$R[1, 1] = \text{norm}(v)$

$Q[:, 1] = v / R[1, 1]$

**for**  $j = 2$  to  $n$  **do**

$v = A[:, j]$

**for**  $i = 1$  to  $j - 1$  **do**

$R[i, j] = Q[:, i]^T A[:, j]$

$v = v - R[i, j] Q[:, i]$

**end for**

$R[j, j] = \text{norm}(v)$

$Q[:, j] = v / R[j, j]$

**end for**



# Singular value decomposition (SVD)

- We would like to work with diagonal matrices, but  $A = S\Lambda S^{-1}$  doesn't produce orthogonal  $S$  in general
- We instead use the following factorization  $A = U\Sigma V^T$ , where  $U$  and  $V$  are orthonormal and  $\Sigma$  is a diagonal matrix
- This is similar to  $Q\Lambda Q^T$ , but the left and right orthogonal matrices are not the same.

- Then,

Instead of eigenvalues  $\Lambda$ , we have singular values  $\Sigma$

Instead of eigenvectors  $S$ , we have left and right singular vectors  $U$  and  $V$

Instead of  $Ax = \lambda x$ , we have  $Av = \sigma u$

Instead of  $AS = S\Lambda$ , we have  $AV = U\Sigma$

# Singular value decomposition (SVD)

- When calculating  $A^\top A$ , we have

$$A^\top A = V \Sigma^\top U^\top U \Sigma V^\top = V \Sigma^\top \Sigma V^\top$$

which is similar to  $Q \Lambda Q^\top$

- The diagonal elements  $\sigma_i^2$  are the positive eigenvalues of  $A^\top A$
- $V$  contains the orthogonal eigenvectors of  $A^\top A$
- $U$  contains the orthogonal eigenvectors of  $AA^\top$
- $U$  and  $V$  provide perfect bases for column and row space of  $A$ ,

$$A_{m \times n} = U_{m \times r} \Sigma_{r \times r} V_{r \times n}^\top$$

# Singular value decomposition (SVD)

- The above presentation is the reduced form.
- By adding any orthonormal basis  $v_{r+1}, \dots, v_n$  from the null space of  $A$ , and any orthonormal basis  $u_{r+1}, \dots, u_m$  from the null space of  $A^\top$ , and completing  $\Sigma$  to an  $m \times n$  matrix by adding zeros, we can write the full SVD as follows:

$$A_{m \times n} = U_{m \times m} \Sigma_{m \times n} V_{n \times n}^\top$$

- Ordering  $\sigma_1 \geq \dots \geq \sigma_r > 0$ , we can write

$$A = u_1 \sigma_1 v_1^\top + \dots + u_r \sigma_r v_r^\top = \sum_{i=1}^r \sigma_i u_i v_i^\top$$

- Left and right singular vectors are also known as Karhunen-Loève bases

# Pseudo inverse

- Using SVD, we can now define a more general concept of inverse
- Note that  $Av_i = \sigma_i u_i$  can be interpreted as  $A$  taking a vector (a basis in this case) from the row space to the column space

- We can define  $A^\dagger$  (called pseudo inverse of  $A$ ) that reverses this operation,

$$A^\dagger u_i = v_i / \sigma_i, \quad i \leq r; \quad A^\dagger u_i = 0, \quad i > r$$

- Singular values of  $A^\dagger$  are  $\Sigma^\dagger = \text{diag}(1/\sigma_1, \dots, 1/\sigma_r)$ ,

$$A^\dagger = V \Sigma^\dagger U^\top$$

- If  $\text{rank}(A) = n$ , then  $A^\dagger = (A^\top A)^{-1} A^\top$  and  $A^\dagger A = I_n$
- If  $\text{rank}(A) = m$ , then  $A^\dagger = A^\top (AA^\top)^{-1}$  and  $AA^\dagger = I_m$

- If  $A$  is square and invertible, then  $A^\dagger = A^{-1}$

# Condition number

- When solving linear systems,  $Ax = b$ , we are interested in measuring the sensitivity of the results to small fluctuations in inputs (e.g., round-off error due to floating-point representation)
- That is, we want to measure  $\Delta x$  given  $\Delta b$
- Suppose  $A$  is positive definite

$$\Delta x = A^{-1} \Delta b$$

# Condition number

- Recall that the eigenvalues of  $A^{-1}$  are  $1/\lambda(A)$
- Therefore,  $1/\lambda_{\min}(A)$  is the largest eigenvalue, and vectors along with the corresponding eigenvectors have the maximum stretch,

$$\|\Delta x\| \leq \|\Delta b\|/\lambda_{\min}(A)$$

- It is better to work with relative errors,

$$\frac{\|\Delta x\|}{\|x\|} \leq \frac{\lambda_{\max}(A)}{\lambda_{\min}(A)} \frac{\|\Delta b\|}{\|b\|}$$

- The term  $c(A) = \frac{\lambda_{\max}(A)}{\lambda_{\min}(A)}$  is called the condition number

# Condition number

- For non-symmetric matrices, the error might blow up along vectors other than eigenvectors.
- In such case, we use the norm  $\|A\|$  instead, which is defined as

$$\max_{x \neq 0} \frac{\|Ax\|}{\|x\|}$$

- Then, the condition number is then defined as  $c(A) = \|A\| \|A^{-1}\|$
- As a rule of thumb, the computer loses  $\log c$  decimals to roundoff error

# Least squares estimation



# Least squares estimates

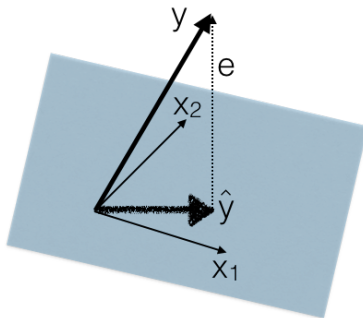
- We now discuss least square estimates for  $X\beta = y$ , where  $X$  is a  $n \times p$  ( $n > p$ ) matrix
- This doesn't have any solution:  $X^{-1}$  doesn't exist; the system is overdetermined (too many equations)
- Instead, we find a solution  $\hat{\beta}$  such that

$$X\hat{\beta} = \hat{y}; \quad y = \hat{y} + e$$

- We find the best solution  $\hat{\beta}$  by making  $e$  small so  $y$  and  $\hat{y}$  are “close” to each other
- We can minimize  $\|e\|^2 = \|y - X\hat{\beta}\|^2 = (y - X\hat{\beta})^\top (y - X\hat{\beta})$

# Least squares estimates

- Geometrically, however,  $e$  would be small when it's perpendicular to  $\hat{y}$  and the column space of  $X$



# Least squares estimates

- Recall that  $N(X^\top) = c(X)^\perp$ ;  $e$  is in the null space of  $X^\top$

$$X^\top e = 0$$

$$X^\top (y - \hat{y}) = 0$$

$$X^\top (y - X\hat{\beta}) = 0$$

- From this, we get the following normal equation,

$$X^\top X\hat{\beta} = X^\top y$$

- Therefore,

$$\hat{\beta} = (X^\top X)^{-1} X^\top y$$

$$\hat{y} = X\hat{\beta} = X(X^\top X)^{-1} X^\top y = Hy$$

# Least squares estimates

- To find  $\hat{\beta}$ , we can solve the normal equation directly, however, this could create some computational difficulties if the condition number of  $(X^T X)$  (which is the square of the condition number of  $X$ ) could be large
- To avoid this, we could use orthogonalization  $X = QR$ , then

$$\begin{aligned}X^T X \hat{\beta} &= X^T y \\(QR)^T QR \hat{\beta} &= (QR)^T y \\R^T R \hat{\beta} &= R^T Q^T y \\R \hat{\beta} &= Q^T y\end{aligned}$$

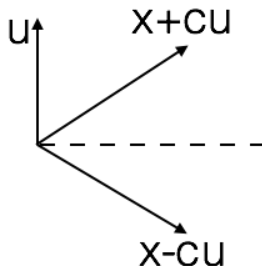
- We find  $Q^T y$ , then use back substitution

# The Householder algorithm

- To find  $QR$ , we could use Gram-Schmidt as discussed before
- Alternatively, we can use the Householder algorithm
- For this, we use reflectors  $H = I - 2uu^\top$ , where  $u$  is a unit length vector
- $H$  is symmetric and orthogonal:  $H^\top H = (I - 2uu^\top)(I - 2uu^\top) = I$
- $H$  reflects  $u$  to  $-u$ :  $Hu = (I - 2uu^\top)u = u - 2u = -u$
- $H$  doesn't change  $x$  vectors perpendicular to  $u$ :  $Hx = x$

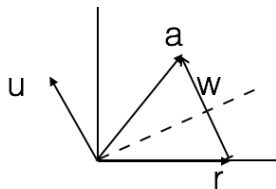
# The Householder algorithm

- In general,  $H$  reflects  $x + cu$  to  $x - cu$ ; the mirror is perpendicular to  $u$



- We can use this fact to create an upper triangular matrix  $R$  by taking each column of  $X$  and creating zeros below its diagonal

# The Householder algorithm



- Consider  $a = \begin{pmatrix} 4 \\ 3 \end{pmatrix}$ , we want to find  $H$  such that  $Ha = r$
- Since  $H$  is orthogonal,  $\|a\| = \|r\|$ , which means  $r = \begin{pmatrix} 5 \\ 0 \end{pmatrix}$
- $w = a - r = \begin{pmatrix} -1 \\ 3 \end{pmatrix}$ , and  $u = w/\|w\| = \frac{1}{\sqrt{10}} \begin{pmatrix} -1 \\ 3 \end{pmatrix}$
- $H = I - 2uu^\top$ , therefore

$$H = \frac{1}{5} \begin{pmatrix} 4 & 3 \\ 3 & -4 \end{pmatrix}$$

# The Householder algorithm

- We continue as above until we find  $H_1, \dots, H_p$
- Applying these reflectors to  $X$  creates  $R$

$$H_p \dots H_1 X = R$$

therefore,

$$X = (H_p \dots H_1)^{-1} R$$

which means

$$Q^{-1} = Q^T = H_p \dots H_1$$

- In practice, we don't need to find  $Q$ , we simply apply the reflectors to  $y$ ,

$$R\hat{\beta} = Q^T y = H_p \dots H_1 y$$

and then use back substitution to find  $\hat{\beta}$



# The Householder algorithm

Initialize  $U_{n \times p} = 0$

**for**  $k = 1$  to  $p$  **do**

$w = X[k:n, k]$

$w[1] = w[1] - \text{norm}(w)$

$u = w / \text{norm}(w)$

$U[k : n, k] = u$

$X[k : n, k : p] = X[k : n, k : p] - 2u(u^\top X[k : n, k : p])$

**end for**

Set  $R_{p \times p}$  to the upper triangular of  $X$

# Least squares using SVD

- Recall the SVD factorization:

$$\begin{aligned}X &= U\Sigma V^{\top} \\ X^{\top}X &= V\Sigma^{\top}\Sigma V^{\top}\end{aligned}$$

- We could use this factorization to solve the normal equation,

$$\begin{aligned}V\Sigma^{\top}\Sigma V^{\top}\hat{\beta} &= V\Sigma^{\top}U^{\top}y \\ V^{\top}\hat{\beta} &= (\Sigma^{\top}\Sigma)^{-1}\Sigma^{\top}U^{\top}y \\ \hat{\beta} &= V\Sigma^{\dagger}U^{\top}y \\ &= X^{\dagger}y\end{aligned}$$

- Here,  $\Sigma^{\dagger} = \text{diag}(1/\sigma_1, \dots, 1/\sigma_n)$  is the pseudoinverse of  $\Sigma$  and  $X^{\dagger}$  is the pseudoinverse of  $X$
- This is the most compact form for least squares estimates

# Recursive least squares

- Suppose we have obtained  $n$  observations,  $X_n$  and  $y_n$ , and found the least squares estimates

$$\hat{\beta}_n = (X_n^\top X_n)^{-1} X_n^\top y_n$$

- Later, we obtain  $k$  more observations, which we denote as  $X_k$  and  $y_k$
- We could of course put all the observations together,  $X_{n+k}$  and  $y_{n+k}$  to obtain the new estimate of regression parameters,  $\hat{\beta}_{n+k}$ , from

$$\begin{aligned} X_{n+k} \beta_{n+k} &= y_{n+k} \\ \begin{pmatrix} X_n \\ X_k \end{pmatrix} \beta_{n+k} &= \begin{pmatrix} y_n \\ y_k \end{pmatrix} \end{aligned}$$

- This would be computationally expensive; instead we can obtain the new estimates iteratively based on the old estimates as follows

# Recursive least squares

- We have

$$\begin{aligned}X_{n+k}^{\top} &= (X_n^{\top} \ X_k^{\top}) \\X_{n+k}^{\top} X_{n+k} &= X_n^{\top} X_n + X_k^{\top} X_k\end{aligned}$$

- The first term is calculated before; next we have

$$X_{n+k}^{\top} y_{n+k} = X_n^{\top} y_n + X_k^{\top} y_k$$

- substituting  $X_n^{\top} X_n$  and multiplying both sides by  $(X_{n+k}^{\top} X_{n+k})^{-1}$ , we have

$$\begin{aligned}\hat{\beta}_{n+k} &= (X_{n+k}^{\top} X_{n+k})^{-1} [(X_{n+k}^{\top} X_{n+k} - X_k^{\top} X_k) \hat{\beta}_n + X_k^{\top} y_k] \\ \hat{\beta}_{n+k} &= \hat{\beta}_n + (X_{n+k}^{\top} X_{n+k})^{-1} X_k^{\top} (y_k - X_k \hat{\beta}_n)\end{aligned}$$

# Recursive least squares

- We can obtain  $R_{n+k}^{-1} = (X_{n+k}^\top X_{n+k})^{-1}$  using the matrix inversion lemma,

$$(A + BCD)^{-1} = A^{-1} - A^{-1}B(C^{-1} + DA^{-1}B)^{-1}DA^{-1}$$

- From the previous slide, we have  $R_{n+k} = X_{n+k}^\top X_{n+k} = (A + BCD)$ , where

$$A = X_n^\top X_n = R_n$$

$$B = X_k^\top$$

$$C = I$$

$$D = X_k$$

- Therefore,

$$R_{n+k}^{-1} = R_n^{-1} - R_n^{-1}X_k^\top(I + X_kR_n^{-1}X_k^\top)^{-1}X_kR_n^{-1}$$

# Weighted least squares

- Recall that we obtained the least squares estimate,  $\hat{\beta}$  by minimizing the length of the residual term  $\|e\| = (e^\top e)^{1/2}$ , which is the Euclidean ( $\ell_2$ ) norm
- In this case, all  $e_i$  elements (all observations) contribute equally
- Sometimes, we want to weight  $e_i$  differently and minimize the weighted sum of residuals instead,  $(e^\top W e)^{1/2}$ , where  $W$  must be positive definite

$$\begin{aligned}X^\top W X \hat{\beta} &= X^\top W y \\ \hat{\beta} &= (X^\top W X)^{-1} X^\top W y\end{aligned}$$

- For example, when  $\text{Cov}(y) = \text{diag}(\sigma_1^2, \dots, \sigma_n^2)$ , we can set

$$W = \text{diag}(1/\sigma_1^2, \dots, 1/\sigma_n^2)$$

# Weighted least squares

- In general, given a positive definite matrix  $P$ ,  $\|x\|_P = (x^\top P x)^{1/2}$  is called the quadratic norm
- Note that this is the same as  $\|P^{1/2}x\|$ , i.e., the Euclidean norm after transformation of  $x$

# Iterative methods

- When solving  $Ax = b$  involves large but sparse matrices, instead of solving the system directly, we can start with an initial guess  $x^{(0)}$ , and improve the solution iteratively
- One such approach is called the Jacobi iteration
- For the first equation, we have

$$\begin{aligned}a_{11}x_1 + \dots + a_{1n}x_n &= b_1 \\x_1 &= \frac{1}{a_{11}}[-(a_{12}x_2 + \dots + a_{1n}x_n)] + \frac{1}{a_{11}}b_1\end{aligned}$$

- In general,

$$x_i = [x_i - \frac{1}{a_{ii}} \sum_{j=1}^n a_{ij}x_j] + \frac{1}{a_{ii}}b_i$$



# Iterative methods

- In the vector form,

$$x = [I - P^{-1}A]x + P^{-1}b$$

where  $P = \text{diag}(a_{11}, \dots, a_{nn})$  is the diagonal part of  $A$

- This gives the recipe, called the Jacobi algorithm, for an iterative approach for finding  $x$

$$x^{(k+1)} = [I - P^{-1}A]x^{(k)} + P^{-1}b$$

- The Gauss-Seidel method is similar to the Jacobi method, but uses the components of new  $x$  (i.e.,  $x^{(k+1)}$ ) as soon as they become available

# Iterative methods

- In general,  $P$  is called the preconditioner matrix
- $P$  should be close to  $A$  but it should be much simpler to work with
- Then, we can come up with an iterative method as follows:

$$Ax = b$$

$$Px = (P - A)x + b$$

$$x = (I - P^{-1}A)x + P^{-1}b$$

$$x^{(k+1)} = (I - P^{-1}A)x^{(k)} + P^{-1}b$$

# Iterative methods

- Note that, the new error,  $e^{(k+1)} = x^{(k+1)} - x$  can be written in terms of the previous error  $e^{(k)} = x^{(k)} - x$ ,

$$e^{(k+1)} = (I - P^{-1}A)e^{(k)} = Me^{(k)}$$

- To ensure the error is shrinking and we are converging to the true solution, we need  $|\lambda(M)| < 1$  for every eigenvalue
- $\max |\lambda(M)|$  is called the spectral radius, which determines the convergence rate